

SPML Assignment 4 - Machine Learning

Clemens Beissel 4547330
Christian Lammers 4578236

August 26, 2018

1 Overall specification

For this machine learning exercise a $n \times m$ world with obstacles was used. The obstacles consisted of walls. Furthermore did this world consist of so called absorbing states, or states with a reward. The reward can be negative or positive. An example for such a world is the figure 1a.

Both algorithms that had to be programmed for this exercise can be applied to such a world. For the purpose of checking the work of the algorithm, a modification was made to the grid world. The utility values of the state are being displayed in their associated state. This makes checking the algorithm much easier.

2 Value iteration algorithm

2.1 Short description of the algorithm

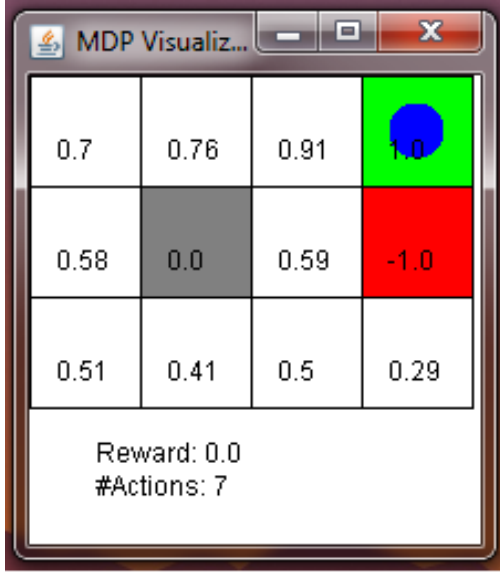
The value iteration algorithm recursively calculates the goodness or the utility of a state s . This calculation is based on the reward of entering that state s and the product of some discount factor γ together with the transition probability p that gives us the probability to enter the successor state s' from state s with action a and the reward for entering the successor state s' .

If the world is non-deterministic performing one action does not always lead to the desired result. Therefore the algorithm has to consider all possible actions that can be applied in state s . For the purpose of calculating the utility of a state, the action with the highest value is chosen.


2.2 Experimenting with different worlds

For every world in this section the values of the states converge. The algorithm finds a policy that leads the agent to the desired goal state.

The agent follows the path with the greatest number, namely from state (1,1) straight to (1,3) and then right to the goal state (4,3) which can also be seen in figure 1a.

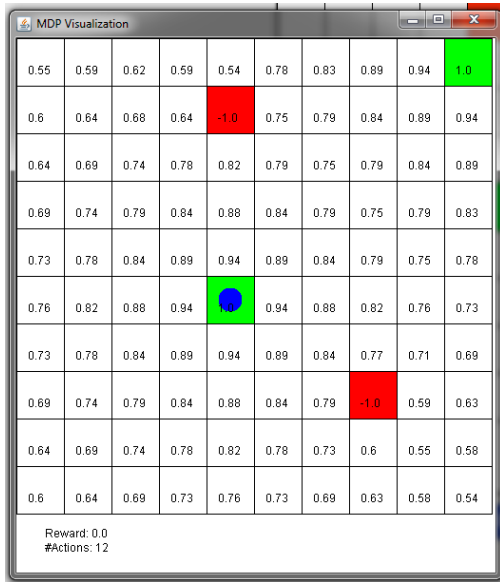


(a) 4x3 world




(b) 10x10 world with winning state at (10,10)

Figure 1: Two different grid worlds




(a) 10x10 world with two negative and positive rewards



(b) 10x10 world with two obstacles

Figure 2: Two different grid worlds

Figure 3: 10x10 world with positive reward at (5,5)

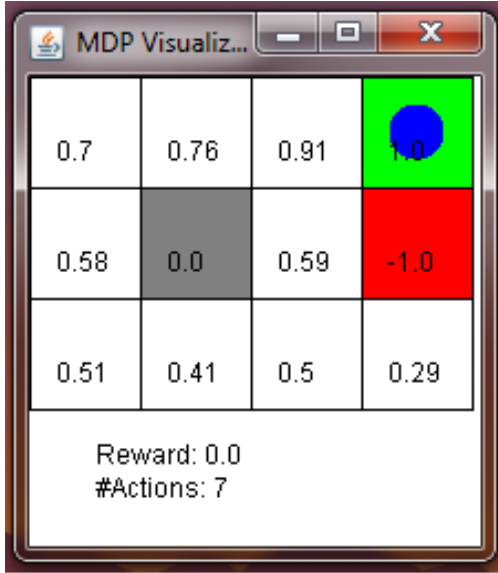


For the other worlds the agent also finds the optimal way to the goal state which is also reflected by the values printed on the state field.

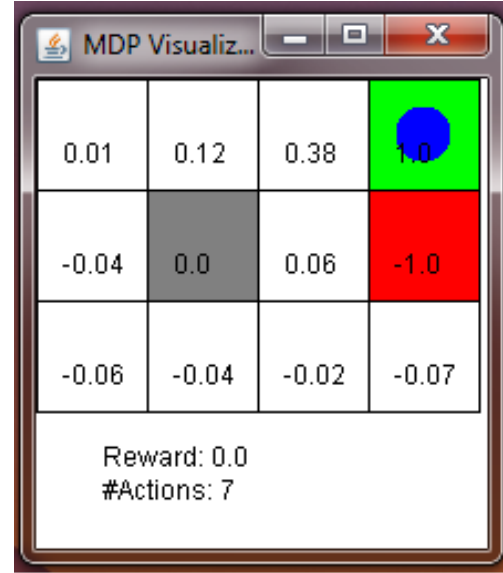
2.3 Experimenting with different values of the discount factor

Experimenting with different values of the discount factor on the 4x3 world with a threshold value of 0.0000001 lead to following results:

After the discount factor was lowered, the difference between the expected values is more subtle. With a discount factor of 0.1 almost all values are hardly any different from each other thus the agent has a hard time following the policy with success. The greater the discount factor is, the closer the expected values are correlated to the values of the absorbing states which can be seen in figure 4a and 4b.



(a) 4x3 world with discount factor 1.0



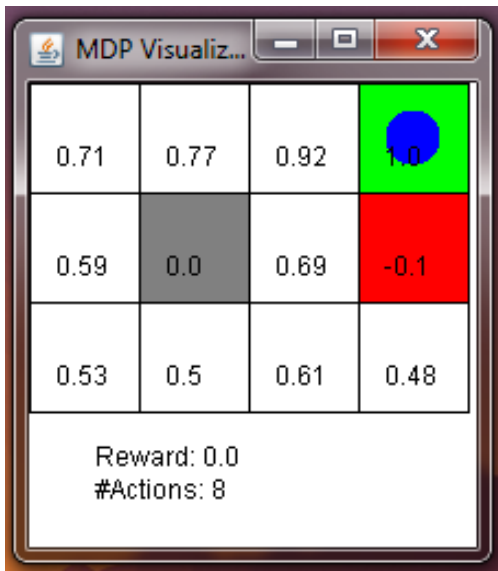
(b) 4x3 world with discount factor 0.5

Figure 4: Two different grid worlds

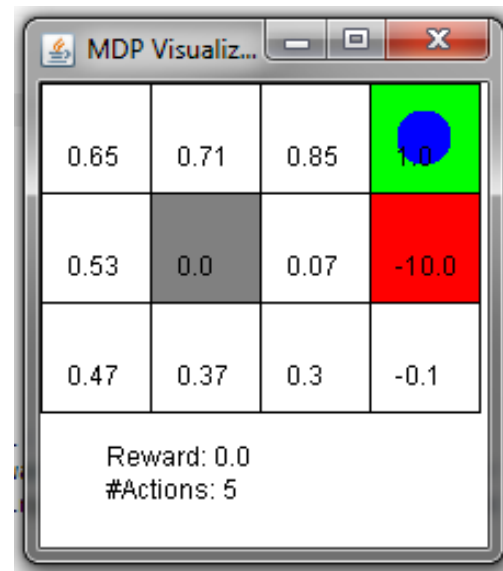
2.4 Experimenting with different values for the state penalty

For the tests with different state penalty values the 4x3 grid world was used with a threshold value of 0.0000001:

As can be seen in figure 5b, do the overall expected values become more negative the lower the negative reward is set (as expected). However the agent does not have any difficulties finding a good path to the positive reward, no matter how negative the expected values become, because there is still the positive absorbing state which is influencing the surrounding states in a positive manner.



(a) 4x3 world with state penalty -0.1



(b) 4x3 world with state penalty -10

Figure 5: Two different grid worlds

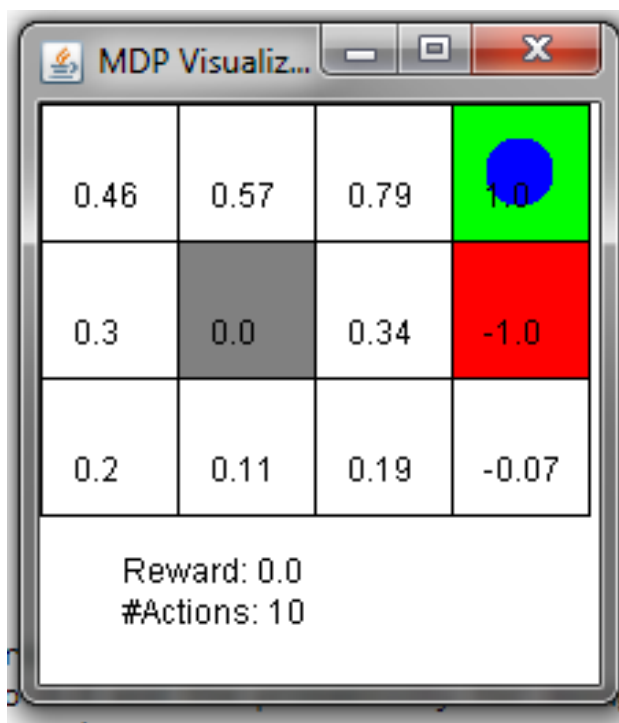
2.5 Experimenting with different settings of the transition probabilities

Normally the probability for moving in the desired direction is set to 0.8 and the probability for doing a sidestep is set to 0.2. Thus moving backwards and staying at the same state is never a possibility. After some experimenting with the transition probabilities new values were established for the backstep and the nostep probability:

```
pPerform = 0.5;
pSidestep = 0.2;
pBackstep = 0.1;
pNoStep = 0.2;
```

The same values as in the previous tests were used for all the other parameters. One can see in figure 6 that the expected values are slightly different from the ones with the original transition probabilities. The overall path, however, does not change due to the unchanged position of the absorbing states. Even after setting the backstep probability to 0.9 the path still does not change.

Figure 6: 4x3 world with different transition probabilities



3 Q-learning algorithm

3.1 Short description of the algorithm

In contrast to the value iteration algorithm, the Q-learning algorithm evaluates the quality of every possible action per state. The agent has to perform actions to receive the corresponding reward and learns through its history of state \rightarrow action \rightarrow reward experiences what actions to take in what state.

This idea was implemented with a three dimensional array where the first two dimensions refer to the location of the state and the third dimension refers to the action. There was a problem with the synchronization of the threads which lead to an incorrect execution of the program when normally running it. This was fixed by inserting a `thread.sleep` in the loop of the algorithm.

3.2 Comparison with exercise 1

Both of the algorithm found a path for the 4x3 grid world that the agent can use to reach the goal state. Both of these paths were found in reasonable time and no significance was found between the performances.

However, comparing the performance of the two algorithms in the 10x10 grid world with each other, shows some obvious differences. The value iteration algorithm converges much faster than the Q-learning algorithm which is mostly due to the actual carrying out of the actions during the Q-learning algorithm. This carrying out of the actions takes much more time and as a consequence does the algorithm need much more time to converge.

A Code used for both algorithms

```
1 package nl.ru.ai.vroon.mdp;
2
3 /**
4  *
5  * @author Clemens Beissel 4547330
6  */
7 public class Tuple {
8
9     private Action a;
10    private Double value;
11
12    public Tuple(Action a, Double v) {
13        this.a = a;
14        value = v;
15    }
16
17    public Action getAction() {
18        return a;
19    }
20
21    public Double getValue() {
22        return value;
23    }
24
25    public void setValue(Double v){
26        value=v;
27    }
28 }
29
30 }
31
32 package nl.ru.ai.vroon.mdp;
33
34 import java.awt.Color;
35 import java.awt.Graphics;
36 import java.awt.Graphics2D;
37 import java.math.BigDecimal;
38 import java.math.RoundingMode;
39
40 import javax.swing.JPanel;
41
42 /**
43  * Creates visual content in accordance with the given MDP
44  * @author Sjoerd Lagarde + some adaptations by Jered Vroon
45  */
```

```

15  */
16  public class DrawPanel extends JPanel {
17
18      private static final long serialVersionUID = 1L;
19      private int screenWidth;
20      private int screenHeight;
21      private MarkovDecisionProblem mdp;
22      private ValueIterationAlgorithm va;
23      private QLearningAlgorithm ql;
24      private boolean valueIteration;
25
26      /**
27       * Constructor
28       * @param mdp
29       * @param screenWidth
30       * @param screenHeight
31       */
32      public DrawPanel(MarkovDecisionProblem mdp, ValueIterationAlgorithm va, int
          screenWidth, int screenHeight) {
33          this.mdp = mdp;
34          this.va = va;
35          this.screenWidth = screenWidth;
36          this.screenHeight = screenHeight;
37          this.valueIteration = true;
38      }
39
40      public DrawPanel(MarkovDecisionProblem mdp, QLearningAlgorithm ql, int
          screenWidth, int screenHeight) {
41          this.mdp = mdp;
42          this.ql = ql;
43          this.screenWidth = screenWidth;
44          this.screenHeight = screenHeight;
45          this.valueIteration = false;
46      }
47
48
49      @Override
50      public void paintComponent(Graphics g) {
51          setBackground(new Color(255, 255, 255));           // White
52          super.paintComponent(g);
53
54          int stepSizeX = screenWidth/mdp.getWidth();
55          int stepSizeY = screenHeight/mdp.getHeight();
56
57          Graphics2D g2 = (Graphics2D)g;
58          for ( int i=0; i<mdp.getWidth(); i++ ) {
59              for ( int j=0; j<mdp.getHeight(); j++ ) {
60                  Field f = mdp.getField(i, j);
61
62                  g2.setPaint(Color.WHITE);
63                  if ( f.equals(Field.REWARD) ) {
64                      g2.setPaint(Color.GREEN);
65                  } else if ( f.equals(Field.NEGREWARD) ) {
66                      g2.setPaint(Color.RED);
67                  } else if ( f.equals(Field.OBSTACLE) ) {
68                      g2.setPaint(Color.GRAY);
69                  }
70                  g2.fillRect(stepSizeX*i, screenHeight - stepSizeY

```

```

71         *(j+1), stepSizeX, stepSizeY);
72
73         if ( mdp.getStateXPosition() == i && mdp.
74             getStateYPostion() == j ) {
75             g2.setPaint(Color.BLUE);
76             g2.fillOval(stepSizeX*i+stepSizeX/4,
77                 screenHeight - stepSizeY*(j+1)+
78                 stepSizeY/4, stepSizeX/2, stepSizeY/2);
79         }
80         g2.setPaint(Color.BLACK);
81         g2.drawRect(stepSizeX*i, screenHeight - stepSizeY
82             *(j+1), stepSizeX, stepSizeY);
83         if (valueIteration)
84             g2.drawString(String.valueOf(round(va.
85                 getExpectedValue(i, j), 2)), stepSizeX*i + 12,
86                 screenHeight - stepSizeY*(j) - 12);
87     }
88 }
89
90     private double round(double value, int places) {
91         if (places < 0) throw new IllegalArgumentException();
92
93         BigDecimal bd = new BigDecimal(value);
94         bd = bd.setScale(places, RoundingMode.HALF_UP);
95         return bd.doubleValue();
96     }
97 }

```

B Code for the value iteration algorithm

```

1  package nl.ru.ai.vroon.mdp;
2
3  import java.util.ArrayList;
4  import java.util.Collection;
5  import java.util.LinkedHashMap;
6  import java.util.List;
7  import java.util.Map;
8  import java.util.Objects;
9  import java.util.stream.Collectors;
10
11  /**
12   *
13   * @author christianlammers
14   */
15  public class ValueIterationAlgorithm {
16
17      MarkovDecisionProblem mdp;
18
19      Double discount = 1.0;

```



```

20     Double sigma = 0.00000001;
21     int counter = 0;
22     private int maxIterations = 10000;
23         Double[][] stateUtilities;
24
25     private Action[][] policy;
26
27     public ValueIterationAlgorithm(MarkovDecisionProblem m) {
28         mdp = m;
29         mdp.setInitialState(0, 0);
30         policy = new Action[mdp.getWidth()][mdp.getHeight()];
31         stateUtilities = new Double[mdp.getWidth()][mdp.getHeight()];
32     }
33
34     /*
35     The core value Iteration algorithm
36     */
37     public void valueIteration() {
38         Map<Action, Double> utils;
39         Double[][] oldStateUtilities = new Double[mdp.getWidth()][mdp.getHeight()];
40         ];
41         initStateUtility(oldStateUtilities);
42         Boolean done = false;
43
44         while (!done || counter >= maxIterations) {
45
46             //loop through every state
47             for (int x = 0; x < mdp.getWidth(); x++) {
48                 for (int y = 0; y < mdp.getHeight(); y++) {
49                     double reward = mdp.getReward(x, y);
50                     if (reward != mdp.getNoReward()) { //if the current state is a
51                         sink or an obstacle
52                         stateUtilities[x][y] = reward;
53                     } else {
54
55                         utils = getUtils(x, y, oldStateUtilities); // returns map
56                         of actions and their values
57                         Tuple bestAction = getBestAction(utils, oldStateUtilities[
58                             x][y]); // returns the best action and the best value
59                         based on utils
60                         stateUtilities[x][y] = reward + bestAction.getValue() *
61                         discount; // gives the current state a new expected
62                         value (bellmann equation)
63
64                         policy[x][y] = bestAction.getAction(); //adds the best
65                         action for every state to the list.
66                     }
67                 }
68             }
69
70             done = checkDifference(stateUtilities, oldStateUtilities, sigma);
71             oldStateUtilities = duplicate(stateUtilities);
72             System.out.println(counter);
73             counter++;
74         }
75     }

```

```

71         for (int x = 0; x < mdp.getWidth(); x++) {
72             for (int y = 0; y < mdp.getHeight(); y++) {
73                 System.out.println("(" + x + "," + y + ")" + " : " +
74                     stateUtilities[x][y] + " Action: " + policy[x][y]);
75             }
76         }
77     }
78 }
79
80 /*
81 returns a HashMap of Actions and their corresponding expected values
82   originating from the position x y.
83 */
84 private Map getUtils(int x, int y, Double[][] oldStateUtilities) {
85     Map<Action, Double> utils = new LinkedHashMap<>();
86     Double thisState = oldStateUtilities[x][y];
87     int right = x + 1;
88     int left = x - 1;
89     int up = y + 1;
90     int down = y - 1;
91
92     if (right < mdp.getWidth()) {
93         utils.put(Action.RIGHT, oldStateUtilities[right][y]);
94     } else {
95         utils.put(Action.RIGHT, thisState);
96     }
97
98     if (left >= 0) {
99         utils.put(Action.LEFT, oldStateUtilities[left][y]);
100     } else {
101         utils.put(Action.LEFT, thisState);
102     }
103
104     if (up < mdp.getHeight()) {
105         utils.put(Action.UP, oldStateUtilities[x][up]);
106     } else {
107         utils.put(Action.UP, thisState);
108     }
109
110     if (down >= 0) {
111         utils.put(Action.DOWN, oldStateUtilities[x][down]);
112     } else {
113         utils.put(Action.DOWN, thisState);
114     }
115
116     return utils;
117 }
118
119 /**
120  * Sets all expected values to 0.0
121  *
122  * @param stateValues
123  */
124 private void initStateUtility(Double[][] oldStateUtilities) {
125     for (int i = 0; i < mdp.getWidth(); i++) {
126         for (int j = 0; j < mdp.getHeight(); j++) {
127             oldStateUtilities[i][j] = 0.0;

```

```

128     }
129 }
130 }
131
132 /*
133 returns the best Action and the corresponding expected value for a state.
134
135 */
136 private Tuple getBestAction(Map<Action, Double> utils, Double valueOfThisState
137 ) {
138     Double[] actionValues = new Double[4];
139
140     actionValues[0] = utils.get(Action.UP) * mdp.getpPerform()
141         + utils.get(Action.nextAction(Action.UP)) * mdp.getPSideStep() / 2
142         + utils.get(Action.previousAction(Action.UP)) * mdp.getPSideStep()
143         + utils.get(Action.backAction(Action.UP)) * mdp.getpBackstep()
144         + valueOfThisState * mdp.getPNoStep();
145     actionValues[1] = utils.get(Action.RIGHT) * mdp.getpPerform()
146         + utils.get(Action.nextAction(Action.RIGHT)) * mdp.getPSideStep()
147         + utils.get(Action.previousAction(Action.RIGHT)) * mdp.
148         getPSideStep() / 2
149         + utils.get(Action.backAction(Action.RIGHT)) * mdp.getpBackstep()
150         + valueOfThisState * mdp.getPNoStep();
151     actionValues[2] = utils.get(Action.DOWN) * mdp.getpPerform()
152         + utils.get(Action.nextAction(Action.DOWN)) * mdp.getPSideStep() /
153         2
154         + utils.get(Action.previousAction(Action.DOWN)) * mdp.getPSideStep
155         () / 2
156         + utils.get(Action.backAction(Action.DOWN)) * mdp.getpBackstep()
157         + valueOfThisState * mdp.getPNoStep();
158     actionValues[3] = utils.get(Action.LEFT) * mdp.getpPerform()
159         + utils.get(Action.nextAction(Action.LEFT)) * mdp.getPSideStep() /
160         2
161         + utils.get(Action.previousAction(Action.LEFT)) * mdp.getPSideStep
162         () / 2
163         + utils.get(Action.backAction(Action.LEFT)) * mdp.getpBackstep()
164         + valueOfThisState * mdp.getPNoStep();
165
166     Double bestValue = -999999999999.0;
167     int bestIndex = 0;
168     for (int i = 0; i < actionValues.length; i++) {
169         if (actionValues[i] > bestValue) {
170             bestValue = actionValues[i];
171             bestIndex = i;
172         }
173     }
174
175     Action bestAction = null;
176     switch (bestIndex) {
177         case 0:
178             bestAction = Action.UP;
179             break;
180         case 1:
181             bestAction = Action.RIGHT;
182             break;
183         case 2:
184             bestAction = Action.DOWN;
185             break;
186         case 3:
187             bestAction = Action.LEFT;
188             break;
189     }
190 }

```

```

179         break;
180     case 3:
181         bestAction = Action.LEFT;
182     }
183
184     return new Tuple(bestAction, bestValue);
185 }
186
187 /*
188 returns true if the difference between the current expected value and the
189 previous expected value is less than sigma for all states.
190 */
191 private Boolean checkDifference(Double[][] stateUtilities, Double[][]
192     oldStateUtilities, Double sigma) {
193     Double difference = 0.0;
194     for (int x = 0; x < mdp.getWidth(); x++) {
195         for (int y = 0; y < mdp.getHeight(); y++) {
196             difference = stateUtilities[x][y] - oldStateUtilities[x][y];
197             if (Math.abs(difference) > sigma) {
198                 return false;
199             }
200         }
201     }
202     return true;
203 }
204
205 /*
206 deep copy of utility array
207 */
208 private Double[][] duplicate(Double[][] stateUtilities) {
209     Double[][] clone = new Double[stateUtilities.length][stateUtilities[0].
210         length];
211     for (int i = 0; i < stateUtilities.length; i++) {
212         for (int j = 0; j < stateUtilities[0].length; j++) {
213             clone[i][j] = stateUtilities[i][j];
214         }
215     }
216     return clone;
217 }
218
219 public Action getPolicy(int x, int y) {
220     return policy[x][y];
221 }
222
223 public double getExpectedValue(int x, int y){
224     return stateUtilities[x][y] == null ? 0.0 : stateUtilities[x][y];
225 }
226
227 /**
228 * the agent performs the moves following the computed policy.
229 */
230 public void ApplyPolicy() {
231     mdp.setShowProgress(true);
232     mdp.setWaittime(500);
233     int xPos = 0;
234     int yPos = 0;
235     while (true) {
236         xPos = mdp.getStateXPosition();

```

```

235         yPos = mdp.getStateYPostion();
236         if (getPolicy(xPos, yPos) != null) {
237             mdp.performAction(getPolicy(xPos, yPos));
238         } else {
239             break;
240         }
241     }
242     System.out.println("finished");
243 }
244
245 }

```

C Code for the Q-Learning algorithm

```

1  package nl.ru.ai.vroon.mdp;
2
3  import java.util.ArrayList;
4  import java.util.Arrays;
5  import java.util.Random;
6
7  /**
8   *
9   * @author christianlammers
10  */
11  public class QLearningAlgorithm {
12
13      MarkovDecisionProblem mdp;
14      double[][][] QValues;
15      int counter = 0;
16      private int maxIterations = 10000;
17      private Random rnd;
18
19      double discount = 0.9;
20      double alpha = 0.2;
21
22      double eps = 1.0;
23      double eps_min = 0.01;
24      double decay = 0.005;
25
26      public QLearningAlgorithm(MarkovDecisionProblem m) {
27          mdp = m;
28          mdp.setInitialState(0, 0);
29          rnd = new Random();
30      }
31
32      public void QLearning() {
33          initializeQTable();
34
35          for (int iter = 0; iter < maxIterations; iter++){
36
37              int x = mdp.getStateXPosition();
38              int y = mdp.getStateYPostion();
39              while (mdp.getField(x, y) != Field.REWARD){
40                  try
41                  {Thread.sleep(0);}
42                  catch (Exception e)
43                  {e.printStackTrace();}
44

```

```

45
46
47
48
49         int a = bestAction(x,y);
50         double expFutValue = actionValue(a);
51         double q = QValues[x][y][a];
52         QValues[x][y][a] = q + alpha*(expFutValue - q);
53
54         x = mdp.getStateXPosition();
55         y = mdp.getStateYPostion();
56
57         eps += -decay;
58         if (eps < eps_min)
59             eps = eps_min;
60     }
61     mdp.restart();
62 }
63
64 printTable();
65
66 mdp.setShowProgress(true);
67 mdp.setWaittime(500);
68 applyPolicy();
69
70
71
72 }
73
74 private void initializeQTable() {
75     QValues = new double[mdp.getWidth()][mdp.getHeight()][4];
76     for (int x = 0; x < mdp.getWidth(); x++) {
77         for (int y = 0; y < mdp.getHeight(); y++) {
78             for (int z = 0; z < QValues[0][0].length; z++)
79                 QValues[x][y][z] = 0.0;
80         }
81     }
82 }
83
84 }
85
86
87 private int bestAction(int x, int y) {
88     float exploitation = rnd.nextFloat();
89
90     int[] possibleActions = getPossibleActions(x,y);
91     int k = 0;
92
93     if (exploitation > eps){ //exploitation{
94
95         //best action
96         double bestValue = -999999999999.0;
97         for (int a : possibleActions){
98             if (QValues[x][y][a] > bestValue){
99                 bestValue = QValues[x][y][a];
100                 k = a;
101             }
102         }
103     }
104     return k;

```

```

104     }
105     else { //exploration
106         k = rnd.nextInt(possibleActions.length);
107         return possibleActions[k];
108     }
109
110     //random action
111     //possibleActions[k];
112 }
113
114 private double actionValue(int a) {
115     double reward = 0;
116     switch(a){
117         case 0: reward = mdp.performAction(Action.UP);
118         break;
119         case 1: reward = mdp.performAction(Action.RIGHT);
120         break;
121         case 2: reward = mdp.performAction(Action.DOWN);
122         break;
123         case 3: reward = mdp.performAction(Action.LEFT);
124     }
125
126     int newX = mdp.getStateXPosition();
127     int newY = mdp.getStateYPostion();
128
129     double maxAction = Arrays.stream(QValues[newX][newY]).max().getAsDouble();
130     //getBestFutureAction(newX, newY);
131     //System.out.println("x: " + newX + " y: " + newY + " value: " + maxAction
132     );
133
134     return reward + discount*maxAction;
135 }
136
137 //public double getExpectedValue(int x, int y)
138
139 private void printTable() {
140     for (int i = 0; i < QValues.length; i++){
141         for (int j =0; j < QValues[0].length; j++){
142             for (int z = 0; z < QValues[0][0].length; z++){
143                 System.out.println("Value for " + "x: " + i + " y: " + j + "
144                 action: " + z + " : " + QValues[i][j][z]);
145             }
146         }
147     }
148 }
149
150 private int[] getPossibleActions(int x, int y) {
151     ArrayList<Integer> possibleActions = new ArrayList<>();
152     for (int a = 0; a < QValues[0][0].length; a++){
153         if (isPossible(x,y,a))
154             possibleActions.add(a);
155     }
156     return possibleActions.stream().mapToInt(i->i).toArray();
157 }
158
159 private boolean isPossible(int x,int y,int a) {
160     switch(a){
161         case 0: return y+1 < mdp.getHeight() && mdp.getField(x, y+1) != Field.
162             OBSTACLE;

```

```

159         case 1: return x+1 < mdp.getWidth() && mdp.getField(x+1, y) != Field.
                OBSTACLE;
160         case 2: return y-1 >= 0 && mdp.getField(x, y-1) != Field.OBSTACLE;
161         case 3: return x-1 >= 0 && mdp.getField(x-1, y) != Field.OBSTACLE;
162     }
163     return true;
164 }
165
166 private double getBestFutureAction(int x, int y) {
167     double[] actionValues = new double[4];
168     actionValues[0] = QValues[x][y][0] * mdp.getpPerform() + QValues[x][y][1]
        * mdp.getPSideStep()/2 + QValues[x][y][2] * mdp.getpBackstep() +
        QValues[x][y][3] * mdp.getPSideStep()/2;
169     actionValues[1] = QValues[x][y][1] * mdp.getpPerform() + QValues[x][y][0]
        * mdp.getPSideStep()/2 + QValues[x][y][2] * mdp.getPSideStep()/2 +
        QValues[x][y][3] * mdp.getpBackstep();
170     actionValues[2] = QValues[x][y][2] * mdp.getpPerform() + QValues[x][y][1]
        * mdp.getPSideStep()/2 + QValues[x][y][3] * mdp.getPSideStep()/2 +
        QValues[x][y][0] * mdp.getpBackstep();
171     actionValues[3] = QValues[x][y][3] * mdp.getpPerform() + QValues[x][y][0]
        * mdp.getPSideStep()/2 + QValues[x][y][2] * mdp.getPSideStep()/2 +
        QValues[x][y][1] * mdp.getpBackstep();
172
173     return Arrays.stream(actionValues).max().getAsDouble();
174 }
175
176 private void applyPolicy() {
177     int x = mdp.getStateXPosition();
178     int y = mdp.getStateYPostion();
179     while (mdp.getField(x, y) != Field.REWARD && mdp.getField(x, y) != Field.
        NEGREWARD){
180         int bestAction = 0;
181         double bestValue = -9999999;
182         int[] possibleActions = getPossibleActions(x,y);
183         for (int a : possibleActions){
184             if (QValues[x][y][a] > bestValue){
185                 bestValue = QValues[x][y][a];
186                 bestAction = a;
187             }
188         }
189         switch (bestAction){
190             case 0: mdp.performAction(Action.UP);
191                     break;
192             case 1: mdp.performAction(Action.RIGHT);
193                     break;
194             case 2: mdp.performAction(Action.DOWN);
195                     break;
196             case 3: mdp.performAction(Action.LEFT);
197                     break;
198         }
199
200         x = mdp.getStateXPosition();
201         y = mdp.getStateYPostion();
202     }
203 }
204 }

```